
CI+SIL × zz_Workflow

Rationale for Coupling Continuous Integration with the LLM-Assisted Prototype Engine

David Ramírez-Betancourt

Universidad Nacional de Colombia — GCPDS

October 2026

Scope. This note justifies, from the literature and from the project's own context documents, why Continuous Integration with Software-in-the-Loop (CI+SIL) is the first integration tier for the zz_Workflow engine in the SDR FM sensor project.

1 Introduction

The SDR FM sensor project uses an LLM-assisted workflow, `zz_Workflow`, to produce verified prototypes. The workflow constrains a language model inside four context-bounded stages with human-gated transitions, and it ends in a prototype accompanied by recorded evidence.

What the project still lacks is a robust strategy for *continuous* verification: once a prototype exists, every later change to code, prompt, model, or context document can silently invalidate the evidence that justified it. LLM-assisted systems make this risk sharper, because the generator is probabilistic and its behaviour can shift without any change to the source tree [17, 3].

This note compares the main validation methodologies against what `zz_Workflow` already produces, and argues for one coupling. **Thesis:** CI+SIL is the natural HIGH and clean coupling. The regression suites that `zz_Workflow` already requires are, in effect, a CI harness, and replaying controlled I/Q in software needs no live RF. The coupling preserves the engine’s central rule that the model never approves its own work. It also keeps claims honest: it delivers *verification*, and it leaves live *validation* to downstream tiers.

Sections 2 to 4 give the foundation and the methodological landscape. Sections 5 and 6 cover AI-assisted hypothesis validation and the engine itself. Sections 7 and 8 analyse and then refine the coupling, and Section 9 concludes.

2 Foundation: Current AI Agents

Probabilistic, not deterministic. LLM outputs vary even at nominal zero temperature; the *background temperature* induced by the inference stack makes identical prompts yield different results [17]. Any workflow that assumes repeatability must therefore measure it rather than assume it.

A degradation ceiling. Agent reliability decays roughly geometrically with the number of sequential steps, with a practical ceiling near 16 steps [3]. If each step succeeds with probability p , a chain of n steps succeeds with p^n ; the projected success at about 100 steps is only ≈ 0.24 . Long autonomous chains are structurally fragile, which argues for short, bounded stages.

Incentives favour hallucination. Benchmarks that score accuracy alone reward confident guessing over abstention, so models are trained to hallucinate [2]. Self-reported correctness is therefore weak evidence.

Structure beats instructions. Reliability gains come mainly from workflow structure (typed interfaces, gates, retries, explicit state) rather than from longer instructions [5]. Likewise, **context engineering outranks prompt engineering**: what the model sees at each step, and nothing more, is the controllable variable [4].

Orchestration under constraints. Surveyed agent systems orchestrate multiple tools while managing state, cost, and safety constraints [1]. These are the same constraints a verified-prototype pipeline must make explicit.

3 Hardware Control via Tools

An LLM only produces text. It reaches hardware through *tools*, using function calling and the Model Context Protocol (MCP) [6, 1]. The resulting closed loop is: agent $\rightarrow$ tool/MCP $\rightarrow$ actuator $\rightarrow$ sensor $\rightarrow$ feedback to the agent.

Two-level control. The LLM plans at the slow level; a deterministic controller executes the fast loop. The model never sits inside a real-time loop. **Physical limits** declared in the tool’s JSON Schema (ranges, units, enumerations) are the first line of defence, since a call that violates the schema is rejected before it reaches the device. Schema clarity also correlates with tool-call accuracy [1].

Two human barriers complement this. *Human-in-the-loop* (HITL) is an **authorization** barrier: a person approves consequential actions. *Hardware-in-the-loop* (HIL) is a **validation** barrier: the control software is exercised against real or emulated hardware [7]. Figure 1 shows both.




Figure 1: Two-level agent–hardware control loop. The LLM plans; schema-bounded tools and a deterministic controller act; HITL authorizes, HIL validates.

4 Validation Methodologies Compared

Five approaches are in common use: the **V-model** [10, 16], **CI+SIL with TDD**, **HIL benches** [7, 11], **LLM-as-judge with HITL** [14, 13], and **multi-criteria decision analysis (MCDA)** [15].

They share four principles: (i) verification is not validation; (ii) checks run from cheap to expensive; (iii) an *external oracle* decides; and (iv) humans act at the extremes (framing and final decision). In short, *nobody validates in one step; you validate in layers*.

The need is empirical. Roughly 83% of autonomous-research agents publish code, but only 38% provide reproducible artifacts [8]. And the key structural rule is that the **oracle must not be the generator** [9].

Figure 2 maps the layers to six phases and marks what zz_Workflow currently covers.




Figure 2: Six-phase validation flow with gates G0–G4 and feedback paths. Phases 1, 3, 5 are covered by zz_Workflow; phase 2 is partial; phases 4 and 6 are gaps.

5 Validating Hypotheses with AI

Code is AI’s strongest contribution. The generate $\rightarrow$ execute $\rightarrow$ test $\rightarrow$ repair loop turns model output into something an external executor can accept or reject. Heinrichs et al. report hundreds of simulate–evaluate–modify iterations with an LLM scripting CST, ADS, and KiCad [12].

LLM-as-judge needs calibration. Agreement with human raters should reach Cohen’s $\kappa \geq 0.6$ (and ≥ 0.8 for strong agreement), with explicit controls for position, verbosity, and self-preference

bias [14, 9].

Simulation has a ceiling. HIL enters when the restriction is hardware [7]. HITL retains irreducible value: in the reported work, humans detected geometry idealizations and routing inconsistencies that automated evaluation missed [13, 12].

“Validated in simulation” does not mean “validated.” The claim must name the environment in which the evidence was produced.

6 The zz_Workflow Engine

zz_Workflow has four context-bounded stages: **Ctx 00** (Problem Baseline), **Ctx 01** (Requirements), **Ctx 02** (Architecture), and **Ctx 03** (Implementation Contract). Transitions are human-gated: the model drafts but never approves. The governing rule is:

Acceptance is decided by pre-declared criteria applied to recorded measurements, never by model judgement.

The terminal output is a **Verified Prototype + Evidence**, which is explicitly *verification*, not validation. When a check fails, **R0–R3 defect routing** returns the classified failure to the responsible stage.

Six governing artifacts support the engine. The most relevant here is **Artefact 3**: the evaluation and benchmarking suites, which act as regression checks after *any* model, prompt, or context change. Figure 3 summarises the structure.




Figure 3: zz_Workflow stages with human gates (H), R0–R3 defect routing back to the responsible stage, the six governing artifacts, and the terminal output.

7 Coupling Analysis: Why CI+SIL

Table 1 rates each methodology by how directly it couples to what zz_Workflow already produces.

Table 1: Coupling of each validation methodology with zz_Workflow.

Approach	Coupling	Reason
CI+SIL	HIGH / clean	Artefact 3 is already a regression harness; SIL replays I/Q without RF.
V&V (V-model)	HIGH-MEDIUM	Matches the stage structure but adds heavy traceability overhead.
LLM-judge + HITL	MEDIUM	Useful for review, but conflicts with deterministic acceptance.
HIL bench	MEDIUM-LOW	Needs physical infrastructure not yet in the loop.
MCDA	LOW	Applies to multi-hypothesis choice, not to single-prototype verification.

Artefact 3 already is the harness. Coupling CI to it requires no new concept: the suites run automatically instead of manually.

The name is precise. “Verified” ≠ “validated”: the output claims conformance to pre-declared criteria, not fitness in the field.

Omitting LLM-as-judge for acceptance is a strength. When the same component proposes and evaluates, optimization pressure exploits evaluator weaknesses [9]. The recipe is role separation plus deterministic gates, where *mechanical rejection beats LLM approval, never the reverse*.

Against the six phases of Figure 2: phases 1, 3, 5 are **covered**; phase 2 is **partial**; phases 4 and 6 are **gaps**.

8 The Refined Coupling

The coupling operates as follows (Figure 4).

- **Seed:** the evidence from the verified prototype (golden I/Q, Pd/Pfa curves) becomes the regression baseline.
- **Version control:** code and the Ctx 03 contract live together as an executable specification.
- **Trigger:** a change to the model, prompt, or any Ctx document, not only code commits. This is the natural CI adaptation for LLM-assisted systems.
- **SIL:** replay of controlled I/Q (synthetic or recorded); no live RF.
- **Regression:** Pd, Pfa, and latency compared against the evidence baseline.
- **Acceptance gate:** deterministic, using the criteria from Ctx 01.
- **Telemetry** (Artefact 5) records failure classes and points to the Ctx to revise. The R0–R3 routing is reused: CI states *which* stage receives the defect.

Live-RF HIL remains a downstream gap.

9 Conclusion

CI+SIL is the natural first integration tier for zz_Workflow. The coupling is HIGH and clean because Artefact 3 already is the regression harness: the evidence produced at the end of the workflow seeds the baseline, and the Ctx 03 contract together with the code serves as an executable specification under version control.




Figure 4: CI+SIL coupled with zz_Workflow. A model, prompt, or context change triggers SIL replay and regression checks against the evidence baseline; a deterministic gate decides, and failures are routed by R0–R3 to the responsible stage while telemetry records the failure class.

Deterministic gates preserve the principle that the model never approves its own work. Triggering on model, prompt, or context changes, and not only on code commits, fits the real failure modes of LLM-assisted systems. Reusing the R0–R3 routing lets CI say which stage must revise the defect.

Two gaps remain downstream: an HIL bench for live RF validation, and MCDA for selecting among multiple hypotheses. Keeping the distinction between *verified* and *validated* makes the claims honest and the next steps tractable.

References

- [1] Harbin Institute of Technology & Harvard, “The Evolution of Tool Use in LLM Agents,” arXiv:2603.22862, Mar. 2026.
- [2] A. T. Kalai *et al.*, “Evaluating LLMs for accuracy incentivizes hallucinations,” *Nature*, vol. 653, Apr. 2026.
- [3] Mittal, “How Fast Do Agents Rot?,” arXiv:2609.01660, Sep. 2026.
- [4] Anthropic, “Effective context engineering for AI agents,” 2026.
- [5] Nerd Level Tech, “AI Agent Reliability 2026: Structure Beats Instructions,” 2026.
- [6] Model Context Protocol Specification, rev. 2025-06-18, modelcontextprotocol.io.
- [7] OPAL-RT, “Guide to HIL Testing,” updated Sep. 2026.
- [8] “Autonomous Research Agents: Verification Gap,” arXiv:2608.05179, Jun. 2026.
- [9] “LLM-as-a-Judge Is Not an Oracle,” arXiv:2609.02246, Sep. 2026.
- [10] QMSdoc, “Traceability in V-Model,” Jan. 2026.
- [11] Sheridan, “Testing Embedded Systems 2026,” Apr. 2026.
- [12] Heinrichs *et al.*, “From Prompt to Prototype: LLM-Driven RF Workflow,” arXiv:2608.31006, Aug. 2026.

- [13] Sadeqi Bajestani *et al.*, “HITL+LLM in smart manufacturing,” *J. Manufacturing Systems*, vol. 86, Jun. 2026.
- [14] FutureAGI, “LLM-as-a-Judge in 2026,” updated Aug. 2026.
- [15] *Sustainable Futures*, “MCDA reliability,” Jun. 2026.
- [16] Sheridan Tech, “Design Validation vs Verification,” Mar. 2026.
- [17] Pith, “Introducing Background Temperature,” arXiv:2604.22411, 2026.